

## Dataset Overview:

The sinking of the Titanic is one of the most infamous shipwrecks in history.

On April 15, 1912, during her maiden voyage, the widely considered “unsinkable” RMS Titanic sank after colliding with an iceberg. Unfortunately, there weren’t enough lifeboats for everyone onboard, resulting in the death of 1502 out of 2224 passengers and crew.

While there was some element of luck involved in surviving, it seems some groups of people were more likely to survive than others.

Train.csv will contain the details of a subset of the passengers on board (891 to be exact) and importantly, will reveal whether they survived or not, also known as the “ground truth”.

The test.csv dataset contains similar information but does not disclose the “ground truth” for each passenger.

## Key Features:

pclass: A proxy for socio-economic status (SES)

1st = Upper

2nd = Middle

3rd = Lower

age: Age is fractional if less than 1. If the age is estimated, is it in the form of xx.5

sibsp: The dataset defines family relations in this way...

Sibling = brother, sister, stepbrother, stepsister

Spouse = husband, wife (mistresses and fiancés were ignored)

parch: The dataset defines family relations in this way...

Parent = mother, father

Child = daughter, son, stepdaughter, stepson

Some children travelled only with a nanny, therefore parch=0 for them.

Survived: Survived (contains your binary predictions: 1 for survived, 0 for deceased)

**Input :**

**Load the data:**

# Library

```
library(tidyverse)  # collection of best packages
```

```
library(caret)      # machine learning functions
```

```
library(MLmetrics)  # machine learning metrics
```

```
library(car)        # VIF calculation
```

```
library(class)      # k-NN
```

# Dataset

```
train <- read.csv('train.csv', na.strings=c('', 'NA'))
```

```
test <- read.csv('test.csv', na.strings=c('', 'NA'))
```

```
PassengerId <- test$PassengerId
```

```
glimpse(train)
```

```
dim(train)
```

# Data Cleaning

```
anyDuplicated(train)
```

#Checking null values

```
colSums(is.na(train))
```

```
colSums(is.na(test))
```

# Handling null values

#1. `Cabin`

```
unique(train$Cabin)
```

```
train$Cabin <- replace_na(train$Cabin, 'X0')
```

```
test$Cabin <- replace_na(test$Cabin, 'X0')
```

#2. `Age`

```
train$Surname <- sapply(str_split(train$Name, ','), `[`, 1) %>% str_trim()
```

```
temp <- sapply(str_split(train$Name, ','), `[`, 2)
```

```
train$Title <- sapply(str_split(temp, '\\.'), `[`, 1) %>% str_trim()
train <- train %>% select(-Name)
```

```
test$Surname <- sapply(str_split(test$Name, ','), `[`, 1) %>% str_trim()
temp <- sapply(str_split(test$Name, ','), `[`, 2)
test$Title <- sapply(str_split(temp, '\\.'), `[`, 1) %>% str_trim()
test <- test %>% select(-Name)
```

```
unique(train$Title)
unique(test$Title)
```

```
test[test$Title == 'Dona', 'Title'] = 'Mrs'
```

```
age_by_title <- train %>%
  group_by(Title) %>%
  summarise(median = median(Age, na.rm = TRUE))
```

```
train <- merge(train, age_by_title)
train[is.na(train$Age), 'Age'] <- train[is.na(train$Age), 'median']
train <- train %>% select(-median)
```

```
test <- merge(test, age_by_title)
test[is.na(test$Age), 'Age'] <- test[is.na(test$Age), 'median']
test <- test %>% select(-median)
```

#3. `Embarked`

```
table(train$Embarked)
```

```
train$Embarked <- replace_na(train$Embarked, 'S')
test$Embarked <- replace_na(test$Embarked, 'S')
```

#4. `Fare`

```
fare_by_pclass <- train %>%
  group_by(Pclass) %>%
  summarise(median = median(Fare, na.rm = TRUE))
```

```
train <- merge(train, fare_by_pclass)
train[is.na(train$Fare), 'Fare'] <- train[is.na(train$Fare), 'median']
train <- train %>% select(-median)
```

```
test <- merge(test, fare_by_pclass)
test[is.na(test$Fare), 'Fare'] <- test[is.na(test$Fare), 'median']
test <- test %>% select(-median)
```

```
colSums(is.na(train))
colSums(is.na(test))
```

After handling missing values there are no missing values

```
train <- train %>%
  mutate_at(vars(Pclass, Title, Survived, Sex, Cabin, Embarked), as.factor)
```

```
test <- test %>%
  mutate_at(vars(Pclass, Title, Survived, Sex, Cabin, Embarked), as.factor)
```

```
glimpse(train)
```

# Metrics, Validation, and Class Imbalance

```
prop.table(table(train$Survived))
```

# Modeling

```
train <- train %>% select(-c(Cabin))
test <- test %>% select(-c(Cabin))
head(train)
```

## k-Nearest Neighbors

# Data normalization

```
train_scaled <- scale(x = train %>% select(c('Age', 'SibSp', 'Parch', 'Fare')))
test_scaled <- scale(x = test %>% select(c('Age', 'SibSp', 'Parch', 'Fare')),
  center = attr(train_scaled, "scaled:center"),
  scale = attr(train_scaled, "scaled:scale"))
head(train_scaled)
pred_cols <- train_scaled[,c('Age', 'SibSp', 'Parch', 'Fare')]
head(pred_cols)
target_col <- train$Survived
head(target_col)
```

# Set the random seed for reproducibility  
set.seed(42)

```

# Create the training and test indices
train_indices <- createDataPartition(target_col, p = 0.8, list = FALSE)

# Split the data into training and test sets

train_data <- pred_cols[train_indices, ]
train_labels <- target_col[train_indices]
test_data <- pred_cols[-train_indices, ]
test_labels <- target_col[-train_indices]

# View the head of the training data and target
head(train_data)
head(train_target)

# View the head of the test data and target
head(test_data)
head(test_target)

#KNN

knn_with_distance_measure<-
function(train_data,test_data,train_labels,k,distance_measure){
  predicted_labels<-
knn(train=train_data,test=test_data,cl=train_labels,k=k,prob=TRUE,use.all=TRUE)
  return(predicted_labels)
}

#Set the values of k

k_values<-c(3,5,7)

# Vector to store accuracy
accuracies <- vector()

# Apply KNN
for (k in k_values){
  #Apply KNN with Euclidean distance
  euc_pred <- knn_with_distance_measure(train_data ,test_data ,train_labels,k,"euclidean")

  #Manhattan
  man_pred <- knn_with_distance_measure(train_data ,test_data
,train_labels,k,"manhattan")

  #Maximum distance
  max_pred <- knn_with_distance_measure(train_data ,test_data ,train_labels,k,"maximum")

```

```

#Evaluate accuracy
accuracy_euclidean<-sum(euc_pred==test_labels)/length(test_labels)
accuracy_manhattan<-sum(man_pred==test_labels)/length(test_labels)
accuracy_maximum<-sum(max_pred==test_labels)/length(test_labels)

# Store accuracy
accuracies <- c(accuracies,accuracy_euclidean,accuracy_manhattan,accuracy_maximum)

#Print the accuracy for the current k value
cat("Accuracy for k=",k,"\n")
cat("Euclidean Distance:",accuracy_euclidean,"\n")
cat("Manhattan Distance:",accuracy_manhattan,"\n")
cat("Maximum Distance:",accuracy_maximum,"\n")
cat("\n")

}

#Create a data frame for accuracies
accuracy_df<-data.frame(Distance=rep(c("Euclidean","Manhattan","Maximum"),
length(k_values)),K=rep(k_values,each=3),Accuracy=accuracies)

ggplot(accuracy_df,aes(x=K,y=Accuracy,color=Distance,group=Distance))+
  geom_line()+
  geom_point()+
  labs(title="Accuracy of k-NN with Different Distance Measures",
       x="k",
       y="Accuracy",
       color="Distance Measure")+
  theme_minimal()

```

## Output:

1.

```
> dim(train)
Rows: 891
Columns: 12
$ PassengerId <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34,...
$ Survived    <int> 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0,...
$ Pclass      <int> 3, 1, 3, 1, 3, 3, 1, 3, 3, 2, 3, 1, 3, 3, 3, 2, 3, 2, 3, 3, 2, 2, 3, 1, 3, 3, 3, 1, 3, 3, 1, 1, 3, 2, 1, 1, 3, 3, 3, 3, 3, 2,...
$ Name        <chr> "Braund, Mr. Owen Harris", "Cumings, Mrs. John Bradley (Florence Briggs Thayer)", "Heikkinen, Miss. Laina", "Futrelle, Mrs. J...
$ Sex         <chr> "male", "female", "female", "female", "male", "male", "male", "male", "female", "female", "female", "female", "male", "male", "...
$ Age         <dbl> 22, 38, 26, 35, 35, NA, 54, 2, 27, 14, 4, 58, 20, 39, 14, 55, 2, NA, 31, NA, 35, 34, 15, 28, 8, 38, NA, 19, NA, NA, 40, NA, N...
$ SibSp       <int> 1, 1, 0, 1, 0, 0, 0, 3, 0, 1, 1, 0, 0, 1, 0, 0, 4, 0, 1, 0, 0, 0, 0, 3, 1, 0, 3, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 2, 1, 1, 1, 1,...
$ Parch       <int> 0, 0, 0, 0, 0, 0, 1, 2, 0, 1, 0, 0, 5, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 5, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
$ Ticket      <chr> "A/5 21171", "PC 17599", "STON/O2. 3101282", "113803", "373450", "330877", "17463", "349909", "347742", "237736", "PP 9549", ...
$ Fare        <dbl> 7.2500, 71.2833, 7.9250, 53.1000, 8.0500, 8.4583, 51.8625, 21.0750, 11.1333, 30.0708, 16.7000, 26.5500, 8.0500, 31.2750, 7.85...
$ Cabin       <chr> NA, "C85", NA, "C123", NA, NA, "E46", NA, NA, NA, "G6", "C103", NA, NA, NA, NA, NA, NA, NA, NA, "D56", NA, "A6", NA, NA, ...
$ Embarked    <chr> "S", "C", "S", "S", "S", "Q", "S", "S", "S", "C", "S", "S", "S", "S", "S", "S", "Q", "S", "S", "C", "S", "S", "Q", "S", "S", ...
> dim(train)
[1] 891 12
> |
```

2.

```
> #Checking null values
>
> colSums(is.na(train))
PassengerId  Survived    Pclass      Name      Sex      Age      SibSp      Parch      Ticket      Fare      Cabin      Embarked
0            0            0          0          0          0      177          0          0          0          0      687          2
> colSums(is.na(test))
PassengerId  Pclass      Name      Sex      Age      SibSp      Parch      Ticket      Fare      Cabin      Embarked
0            0          0          0          0          0          0          0          1          0      327          0
> |
```

## 3.After removing null values

```
> colSums(is.na(train))
Pclass      Title  PassengerId  Survived      Sex      Age      SibSp      Parch      Ticket      Fare      Cabin      Embarked
0            0            0            0          0          0          0          0          0          0          0          0
Surname
0
> colSums(is.na(test))
Pclass      Title  PassengerId      Sex      Age      SibSp      Parch      Ticket      Fare      Cabin      Embarked      Surname
0            0            0          0          0          0          0          0          0          0          0          0
> |
```

4.

```
> head(train)
Pclass Title  PassengerId Survived  Sex Age SibSp Parch Ticket Fare Cabin Embarked Surname
1     1  Capt       746         0 male   70   1   1 WE/P 5735 71.0000 B22     S      Crosby
2     1   Col       695         0 male   60   0   0 113800 26.5500 X0      S      Weir
3     1   Col       648         1 male   56   0   0 13213 35.5000 A26     C Simonius-Blumer
4     1   Don       31          0 male   40   0   0 PC 17601 27.7208 X0      C      Uruchurtu
5     1   Dr       633         1 male   32   0   0 13214 30.5000 B50     C Stahelin-Maeglin
6     1   Mr       371         1 male   25   1   0 11765 55.4417 E50     C      Harder
> |
```

5.

```
>
> prop.table(table(train$Survived))
      0      1
0.6161616 0.3838384
> |
```

6.

```
> head(train_scaled)
      Age      SibSp      Parch      Fare
[1,] 3.0613503 0.4325504 0.7671990 0.78070266
[2,] 2.3075051 -0.4742788 -0.4734077 -0.11378180
[3,] 2.0059670 -0.4742788 -0.4734077 0.06632249
[4,] 0.7998146 -0.4742788 -0.4734077 -0.09022135
[5,] 0.1967384 -0.4742788 -0.4734077 -0.03429443
[6,] -0.3309533 0.4325504 -0.4734077 0.46761700
> pred_cols <- train_scaled[,c('Age', 'SibSp', 'Parch', 'Fare')]
> head(pred_cols)
      Age      SibSp      Parch      Fare
[1,] 3.0613503 0.4325504 0.7671990 0.78070266
[2,] 2.3075051 -0.4742788 -0.4734077 -0.11378180
[3,] 2.0059670 -0.4742788 -0.4734077 0.06632249
[4,] 0.7998146 -0.4742788 -0.4734077 -0.09022135
[5,] 0.1967384 -0.4742788 -0.4734077 -0.03429443
[6,] -0.3309533 0.4325504 -0.4734077 0.46761700
> target_col <- train$Survived
```

7.

```
> head(train_data)
      Age      SibSp      Parch      Fare
[1,] 3.0613503 0.4325504 0.7671990 0.78070266
[2,] 2.3075051 -0.4742788 -0.4734077 -0.11378180
[3,] 2.0059670 -0.4742788 -0.4734077 0.06632249
[4,] 0.7998146 -0.4742788 -0.4734077 -0.09022135
[5,] -0.3309533 0.4325504 -0.4734077 0.46761700
[6,] 1.4782753 -0.4742788 -0.4734077 -0.12627440
> head(train_labels)
[1] 0 0 1 0 1 1
Levels: 0 1
>
> # View the head of the test data and target
> head(test_data)
      Age      SibSp      Parch      Fare
[1,] 0.1967384 -0.4742788 -0.4734077 -0.03429443
[2,] 3.1367349 -0.4742788 -0.4734077 0.34813440
[3,] 0.6490455 -0.4742788 -0.4734077 -0.64805768
[4,] 1.4028908 0.4325504 -0.4734077 0.14882837
[5,] 0.7998146 -0.4742788 -0.4734077 -0.64805768
[6,] 0.5736610 -0.4742788 0.7671990 -0.05039314
> head(test_labels)
[1] 1 0 0 1 0 0
Levels: 0 1
>
```

8.

```
Accuracy for k= 3
Euclidean Distance: 0.740113
Manhattan Distance: 0.7457627
Maximum Distance: 0.740113
```


```
Accuracy for k= 5
Euclidean Distance: 0.7457627
Manhattan Distance: 0.740113
Maximum Distance: 0.740113
```

```
Accuracy for k= 7
Euclidean Distance: 0.740113
Manhattan Distance: 0.740113
Maximum Distance: 0.7344633
```


```
> |
```



9.



## Discussion and Analysis:



The graph above shows how the accuracy of the knn algorithm with different values of k on using the 3 distance measuring methods (Euclidean, Manhattan, and Maximum Dimension).

It can be seen that the accuracy was the highest when the value of k was 3 in manhattan and maximum distance measure method. In this instance, Manhattan distance had the highest accuracy.

When the value of k was increased to 5, the accuracies of all the distance methods dropped, except euclidean but the Maximum dimension was the highest. For value k=5 the accuracy to Euclidean distance is high.

In the end, when the value of k was set to 7, the accuracy dropped again. However, this time, Manhattan and Euclidean had the same accuracies and maximum dimension had the lowest accuracy.

From this observation we have found that most accurate value for k is 3. Most accurate value for the method is Euclidean and Manhattan .

Based on this findings , it can be said that the highest accuracy for the value of k(3) which is Manhattan distance .